

CS-E-106: Data Modeling

Assignment 1

Instructor: Hakan Gogtas
Submitted by: Saurabh Kulkarni

Due Date: 09/23/2019

Solution 1:

Regression Function: $Y_i = \beta_0 + \epsilon_i$

Thus, sum of squared residuals, $Q = \sum_{i=1}^n (Y_i - \beta_0 - \epsilon_i)^2$

Taking partial derivatives w.r.t. β_0

$$\frac{\partial Q}{\partial \beta_0} = 2 \sum_{i=1}^n (Y_i - \beta_0)(-1) \stackrel{\text{set}}{=} 0$$

$$\therefore \sum_{i=1}^n Y_i = n\beta_0$$

$$\therefore \beta_0 = \frac{\sum_{i=1}^n Y_i}{n}$$

$$\therefore \beta_0 = \bar{Y}$$

Solution 2:

Install R packages:

```
# install.packages("faraway")
# install.packages("ggplot2")
# install.packages("corrgram")
```

Numerical Summary of the dataset:

```
library(faraway)
teengamb$sex <- as.factor(teengamb$sex)
summary(teengamb)
```

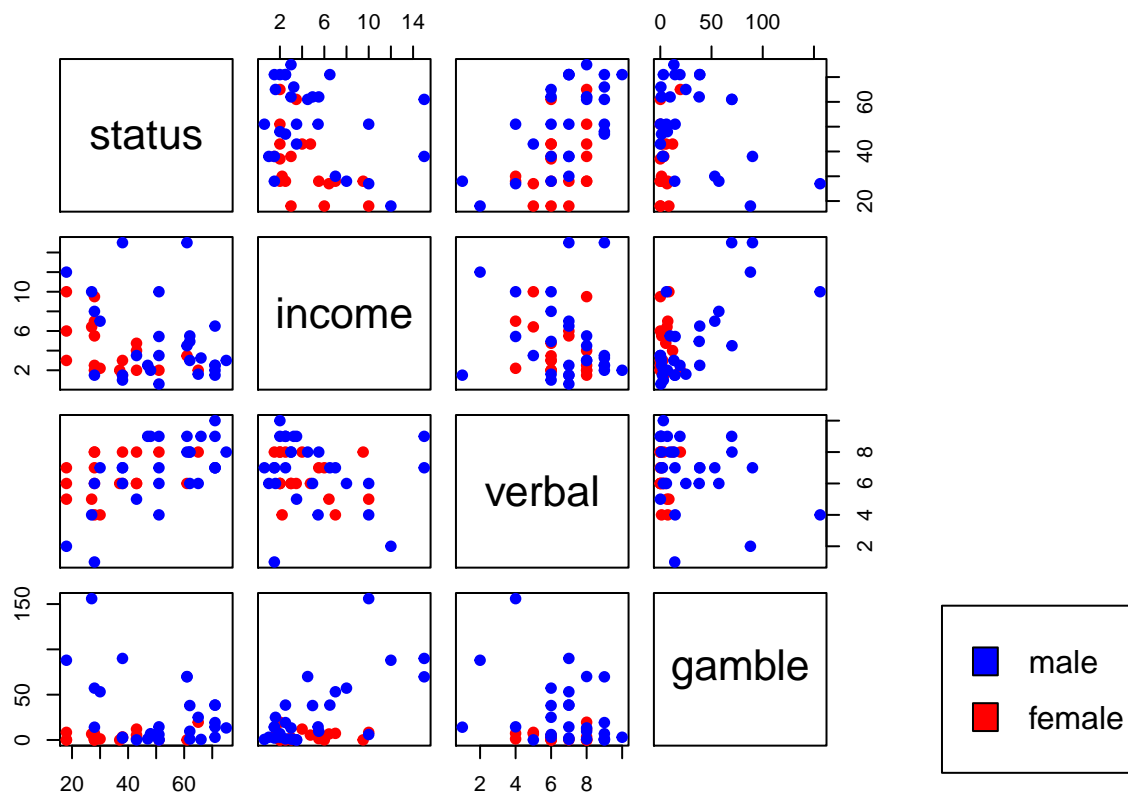
```
##  sex      status      income      verbal      gamble
##  0:28  Min.   :18.00  Min.    : 0.600  Min.    : 1.00  Min.    : 0.0
##  1:19  1st Qu.:28.00  1st Qu.: 2.000  1st Qu.: 6.00  1st Qu.: 1.1
##           Median :43.00  Median : 3.250  Median : 7.00  Median : 6.0
##           Mean   :45.23  Mean    : 4.642  Mean    : 6.66  Mean    :19.3
##           3rd Qu.:61.50  3rd Qu.: 6.210  3rd Qu.: 8.00  3rd Qu.:19.4
##           Max.    :75.00  Max.    :15.000  Max.    :10.00  Max.    :156.0
```

```
help("teengamb")
```

Looking at the summary of the dataset, we can see that the numerical variables are on different scales. `help()` on “teengamb” dataset confirms it.

Scatterplot Matrix of Numerical Variables:

```
pairs(teengamb[, -1], pch=19, col=c('blue', 'red')[teengamb$sex], oma=c(3,3,3,15))
par(xpd = TRUE)
legend("bottomright", fill = c('blue', 'red'), legend = c("male", "female"))
```




We can see that the gambling expenditure increases with the increase in income for males but not a very steep increase for females. We can also see a correlation between the variables `verbal` and `status`.

Correlogram for Males:

```
library(corrgram)

## Registered S3 method overwritten by 'seriation':
##   method      from
##   reorder.hclust gclus

corrgram(teengamb[which(teengamb$sex==0),], order=TRUE, lower.panel=panel.shade,
  upper.panel=panel.pie, text.panel=panel.txt, oma=c(3,3,3,15),
  main="Correlogram for Males")
```



Here, we subset the dataset for males and plot a correlogram to hone in a little bit. The correlogram confirms the correlations we suspected above - we can see strong positive correlation between **gamble** and **income** as well as between **verbal** and **status** for males.

Solution 3:

(a)

reg_loop: Here we write a function to loop over all the desired independent variables in the problem statement, regress the given dependent variable (**Number.of.active.physicians**) on each of them and then return a list of resulting linear models.

```
reg_loop <- function(df, x_cols, y_str) {
  lm_regs = list({})
  for(i in 1:length(x_cols)){
    x_str = x_cols[i]
    formula = as.formula(paste(y_str, x_str))
    lm_regs[[i]] = lm(formula, data=df)
    print(paste("Linear Regression Summary:", x_cols[i]))
    print(summary(lm_regs[[i]]))
  }
  lm_regs
}
```

Fitting the three models:

```
cdi = read.csv("cdi.csv")
cdi = lapply(cdi, as.numeric)
cdi = data.frame(cdi)
x_cols = c("Total.population", "Number.of.hospital.beds", "Total.personal.income")
y_str = "Number.of.active.physicians ~"
lm_fits = reg_loop(df=cdi, x_cols=x_cols, y_str=y_str)
```

```
## [1] "Linear Regression Summary: Total.population"
```

```
##
## Call:
## lm(formula = formula, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1969.4  -209.2   -88.0    27.9   3928.7
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -1.106e+02  3.475e+01  -3.184  0.00156 **
## Total.population  2.795e-03  4.837e-05  57.793  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 610.1 on 438 degrees of freedom
## Multiple R-squared:  0.8841, Adjusted R-squared:  0.8838
## F-statistic: 3340 on 1 and 438 DF, p-value: < 2.2e-16
##
## [1] "Linear Regression Summary: Number.of.hospital.beds"
##
## Call:
## lm(formula = formula, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3133.2  -216.8   -32.0    96.2   3611.1
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -95.93218   31.49396  -3.046  0.00246 **
## Number.of.hospital.beds  0.74312    0.01161  63.995  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 556.9 on 438 degrees of freedom
## Multiple R-squared:  0.9034, Adjusted R-squared:  0.9032
## F-statistic: 4095 on 1 and 438 DF, p-value: < 2.2e-16
##
## [1] "Linear Regression Summary: Total.personal.income"
##
## Call:
## lm(formula = formula, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1926.6  -194.5   -66.6    44.2   3819.0
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -48.39485   31.83333  -1.52   0.129
## Total.personal.income  0.13170    0.00211  62.41  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 569.7 on 438 degrees of freedom
## Multiple R-squared:  0.8989, Adjusted R-squared:  0.8987
## F-statistic: 3895 on 1 and 438 DF,  p-value: < 2.2e-16
```

Thus, we have the *three regression equations* as:

1. Number.of.active.physicians = -110.63 + 0.0028*Total.population
2. Number.of.active.physicians = -95.93 + 0.74*Number.of.hospital.beds
3. Number.of.active.physicians = -48.39 + 0.13*Total.personal.income


(b)


plot_reg_func: Creates three distinct plots for each model created above.

```
plot_reg_func <- function(fits, df, x_cols){
  for(i in 1:length(x_cols)){
    plot(df[[x_cols[i]]], df[["Number.of.active.physicians"]], xlab=x_cols[i], ylab="Number.of.active.physicians",
        abline(fits[[i]]))
  }
}
```

Plots for Fitted Regression Lines:

```
plot_reg_func(fits=lm_fits, df=cdi, x_cols=x_cols)
```





Based on the plots, we can say that the simple linear regression models provide a good fit for each of the three independent variable. We can also confirm this by looking at the R-squared values in the summaries printed above or MSE's in part (c).

(c)

`mse_loop`: This function goes also over the three predictor variables and calculates the MSE for each of them using the respective model from the list we provide to it.

```
mse_loop <- function(fits, df, x_cols){
  for(i in 1:length(x_cols)){
    yHat = predict(fits[[i]], df[x_cols])
    resid = (df$Number.of.active.physicians - yHat)
```

```

    SSE = (sum(resids^2))
    df_resids = (nrow(df)-2)
    MSE = (SSE/df_resids)
    print(paste("MSE for",x_cols[i]))
    print(MSE)
  }
}

```

Calculated MSEs:

```
mse_loop(fits=lm_fits, df=cdi, x_cols=x_cols)
```

```

## [1] "MSE for Total.population"
## [1] 372203.5
## [1] "MSE for Number.of.hospital.beds"
## [1] 310191.9
## [1] "MSE for Total.personal.income"
## [1] 324539.4

```

Thus we can see that `Number.of.hospital.beds` gives the least MSE and thus has least variability around the fitted regression line.

Solution 4:

Here we repeat the whole solution to problem 3 but our models are now trained on the 70% training data. We create this by randomly sampling 70% of the data as shown below and call that as our `train_cdi` and the remaining we call `test_cdi`.

```

train_ind = sample(1:nrow(cdi), 0.7 * nrow(cdi))
test_ind = setdiff(1:nrow(cdi), train_ind)
train_cdi = cdi[train_ind,]
test_cdi = cdi[test_ind,]

```

(a)

Fitting the three models:

```
lm_fits_tr = reg_loop(df=train_cdi, x_cols=x_cols, y_str=y_str)
```

```

## [1] "Linear Regression Summary: Total.population"
##
## Call:
## lm(formula = formula, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1937.7  -214.0  -113.3    15.6   3913.4
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -7.364e+01  4.204e+01  -1.751   0.0809 .
## Total.population  2.763e-03  5.746e-05  48.087  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 628.7 on 306 degrees of freedom
## Multiple R-squared:  0.8831, Adjusted R-squared:  0.8828
## F-statistic: 2312 on 1 and 306 DF, p-value: < 2.2e-16

```

```
##
## [1] "Linear Regression Summary: Number.of.hospital.beds"
##
## Call:
## lm(formula = formula, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3537.1  -214.4   -28.6   110.7  3571.7
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -125.3785     38.9880  -3.216  0.00144 **
## Number.of.hospital.beds    0.7888      0.0149  52.947 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 576.9 on 306 degrees of freedom
## Multiple R-squared:  0.9016, Adjusted R-squared:  0.9013
## F-statistic: 2803 on 1 and 306 DF, p-value: < 2.2e-16
##
## [1] "Linear Regression Summary: Total.personal.income"
##
## Call:
## lm(formula = formula, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1919.4  -203.0   -81.9    27.6   3811.8
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -25.105994     37.408823  -0.671    0.503
## Total.personal.income    0.130432     0.002422  53.855 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 568.1 on 306 degrees of freedom
## Multiple R-squared:  0.9046, Adjusted R-squared:  0.9043
## F-statistic: 2900 on 1 and 306 DF, p-value: < 2.2e-16
```

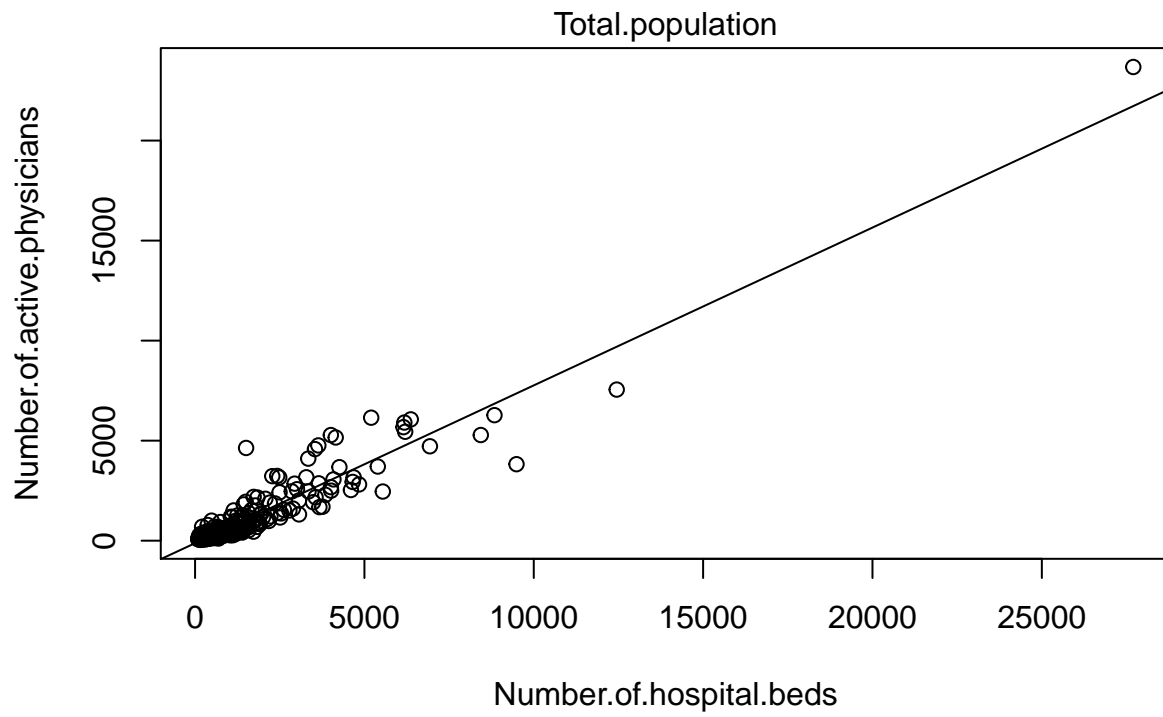

Thus, we have the *new three regression equations* on training set as:


1. $\text{Number.of.active.physicians} = -114.4 + 0.0027 \times \text{Total.population}$
2. $\text{Number.of.active.physicians} = -119.98 + 0.76 \times \text{Number.of.hospital.beds}$
3. $\text{Number.of.active.physicians} = -53.4 + 0.13 \times \text{Total.personal.income}$

(b)

Plots for Fitted Regression Lines:

```
plot_reg_func(fits=lm_fits_tr, df=train_cdi, x_cols=x_cols)
```



The plots more or less show similar fits as in Solution 3.

(c)

Calculated MSEs:

```
mse_loop(fits=lm_fits, df=train_cdi, x_cols=x_cols)
```

```
## [1] "MSE for Total.population"
## [1] 396299.4
## [1] "MSE for Number.of.hospital.beds"
## [1] 344289.6
## [1] "MSE for Total.personal.income"
## [1] 323258.1
```

Thus, with 70% training sample we get `Total.personal.income` as the variable that gives us the smallest variability and thus the best fit.

Solution 5:

(a)

```
lm_gamb = lm(gamble~income, data=teengamb)
summary(lm_gamb)
```

```
##
## Call:
## lm(formula = gamble ~ income, data = teengamb)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -46.020 -11.874  -3.757  11.934  107.120
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -6.325      6.030  -1.049    0.3
```

```
## income          5.520          1.036      5.330 3.05e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 24.95 on 45 degrees of freedom
## Multiple R-squared:  0.387, Adjusted R-squared:  0.3734
## F-statistic: 28.41 on 1 and 45 DF,  p-value: 3.045e-06
```

The regression equation is: $\text{gamble} = -6.325 + 5.52 \cdot \text{income}$

```
yHat = predict(lm_gamb, teengamb["income"])
resids = (teengamb$gamble - yHat)
SSE = (sum(resids^2))
df_resids = (nrow(teengamb) - 2)
MSE = (SSE / df_resids)
print(paste("Mean of the residuals:", mean(resids)))
```

```
## [1] "Mean of the residuals: -2.51361431669039e-15"
```

```
print(paste("Median of the residuals:", median(resids)))
```

```
## [1] "Median of the residuals: -3.7573821062283"
```

```
print(paste("MSE:", MSE))
```

```
## [1] "MSE: 622.41305773688"
```

Note how close the mean of the residuals is to zero. Ideally, it should be perfectly zero (that is one of the properties of least squares).

(b)

```
print(paste("Case number for highest positive residual:", which(resids == max(resids))))
```

```
## [1] "Case number for highest positive residual: 24"
```

```
print(resids[24])
```

```
##          24
```

```
## 107.1197
```

```
print(teengamb[24,])
```

```
##      sex status income verbal gamble
```

```
## 24    0      27      10        4    156
```